

**Lecture 4:**



# **Drawing a Triangle (and an Intro to Sampling)**

---

**Computer Graphics  
CMU 15-462/15-662**

# Assignment 0.5 due, MiniHW 0 due, Assignment 1 out, MiniHW 1 out!

- **GOAL: Implement a basic “rasterizer”**
  - (Topic of today’s lecture)
  - We hand you a bunch of lines, triangles, etc.
  - You draw them by lighting up pixels on the screen!
- **Code skeleton available from course webpage**
- **First checkpoint due February 9th.**
- **Completed assignment due February 16th.**





# Assignment 0.5 due, MiniHW 0 due, Assignment 1 out, MiniHW 1 out!

- **GOAL: Implement a basic “rasterizer”**
  - (Topic of today’s lecture)
  - We hand you a bunch of lines, triangles, etc.
  - You draw them by lighting up pixels on the screen!
- **Code skeleton available from course webpage**
- **First checkpoint due February 9th.**
- **Completed assignment due February 16th.**

*See piazza for  
build tutorial!*



*Extra office  
hours today  
3-5pm for build  
questions!*





# TODAY: Rasterization


- Two major techniques for “getting stuff on the screen”
- Rasterization (TODAY)
  - for each primitive (e.g., triangle), which pixels light up?
  - extremely fast (BILLIONS of triangles per second on GPU)
  - harder (but not impossible) to achieve photorealism
  - perfect match for 2D vector art, fonts, quick 3D preview, ...
- Ray tracing (LATER)
  - for each pixel, which primitives are seen?
  - easier to get photorealism
  - generally slower
  - much more later in the semester!



# 3D Image Generation Pipeline(s)

- Can talk about image generation in terms of a “pipeline”:
  - **INPUTS** — what image do we want to draw?
  - **STAGES** — sequence of transformations from input → output
  - **OUTPUTS** — the final image

E.g., our pipeline from the first lecture:



# Rasterization Pipeline

- Modern real time image generation based on rasterization
  - INPUT: 3D “primitives”—essentially all triangles!
    - possibly with additional attributes (e.g., color)
  - OUTPUT: bitmap image (possibly w/ depth, alpha, ...)
- Our goal: understand the stages in between\*

## INPUT (TRIANGLES)

### VERTICES

A: ( 1, 1, 1 )	E: ( 1, 1, -1 )
B: (-1, 1, 1 )	F: (-1, 1, -1 )
C: ( 1, -1, 1 )	G: ( 1, -1, -1 )
D: (-1, -1, 1 )	H: (-1, -1, -1 )


### TRIANGLES

EHF, GFH, FGB, CBG,  
GHC, DCH, ABD, CDB,  
HED, ADE, EFA, BAF

## RASTERIZATION PIPELINE



## OUTPUT (BITMAP IMAGE)




**\*In practice, usually executed by graphics processing unit (GPU)**





# Why triangles?

- Rasterization pipeline converts all primitives to triangles
  - even points and lines!
- Why?
  - can approximate any shape
  - always planar, well-defined normal
  - easy to interpolate data at corners
    - “barycentric coordinates”
- Key reason: once everything is reduced to triangles, can focus on making an extremely well-optimized pipeline for drawing them



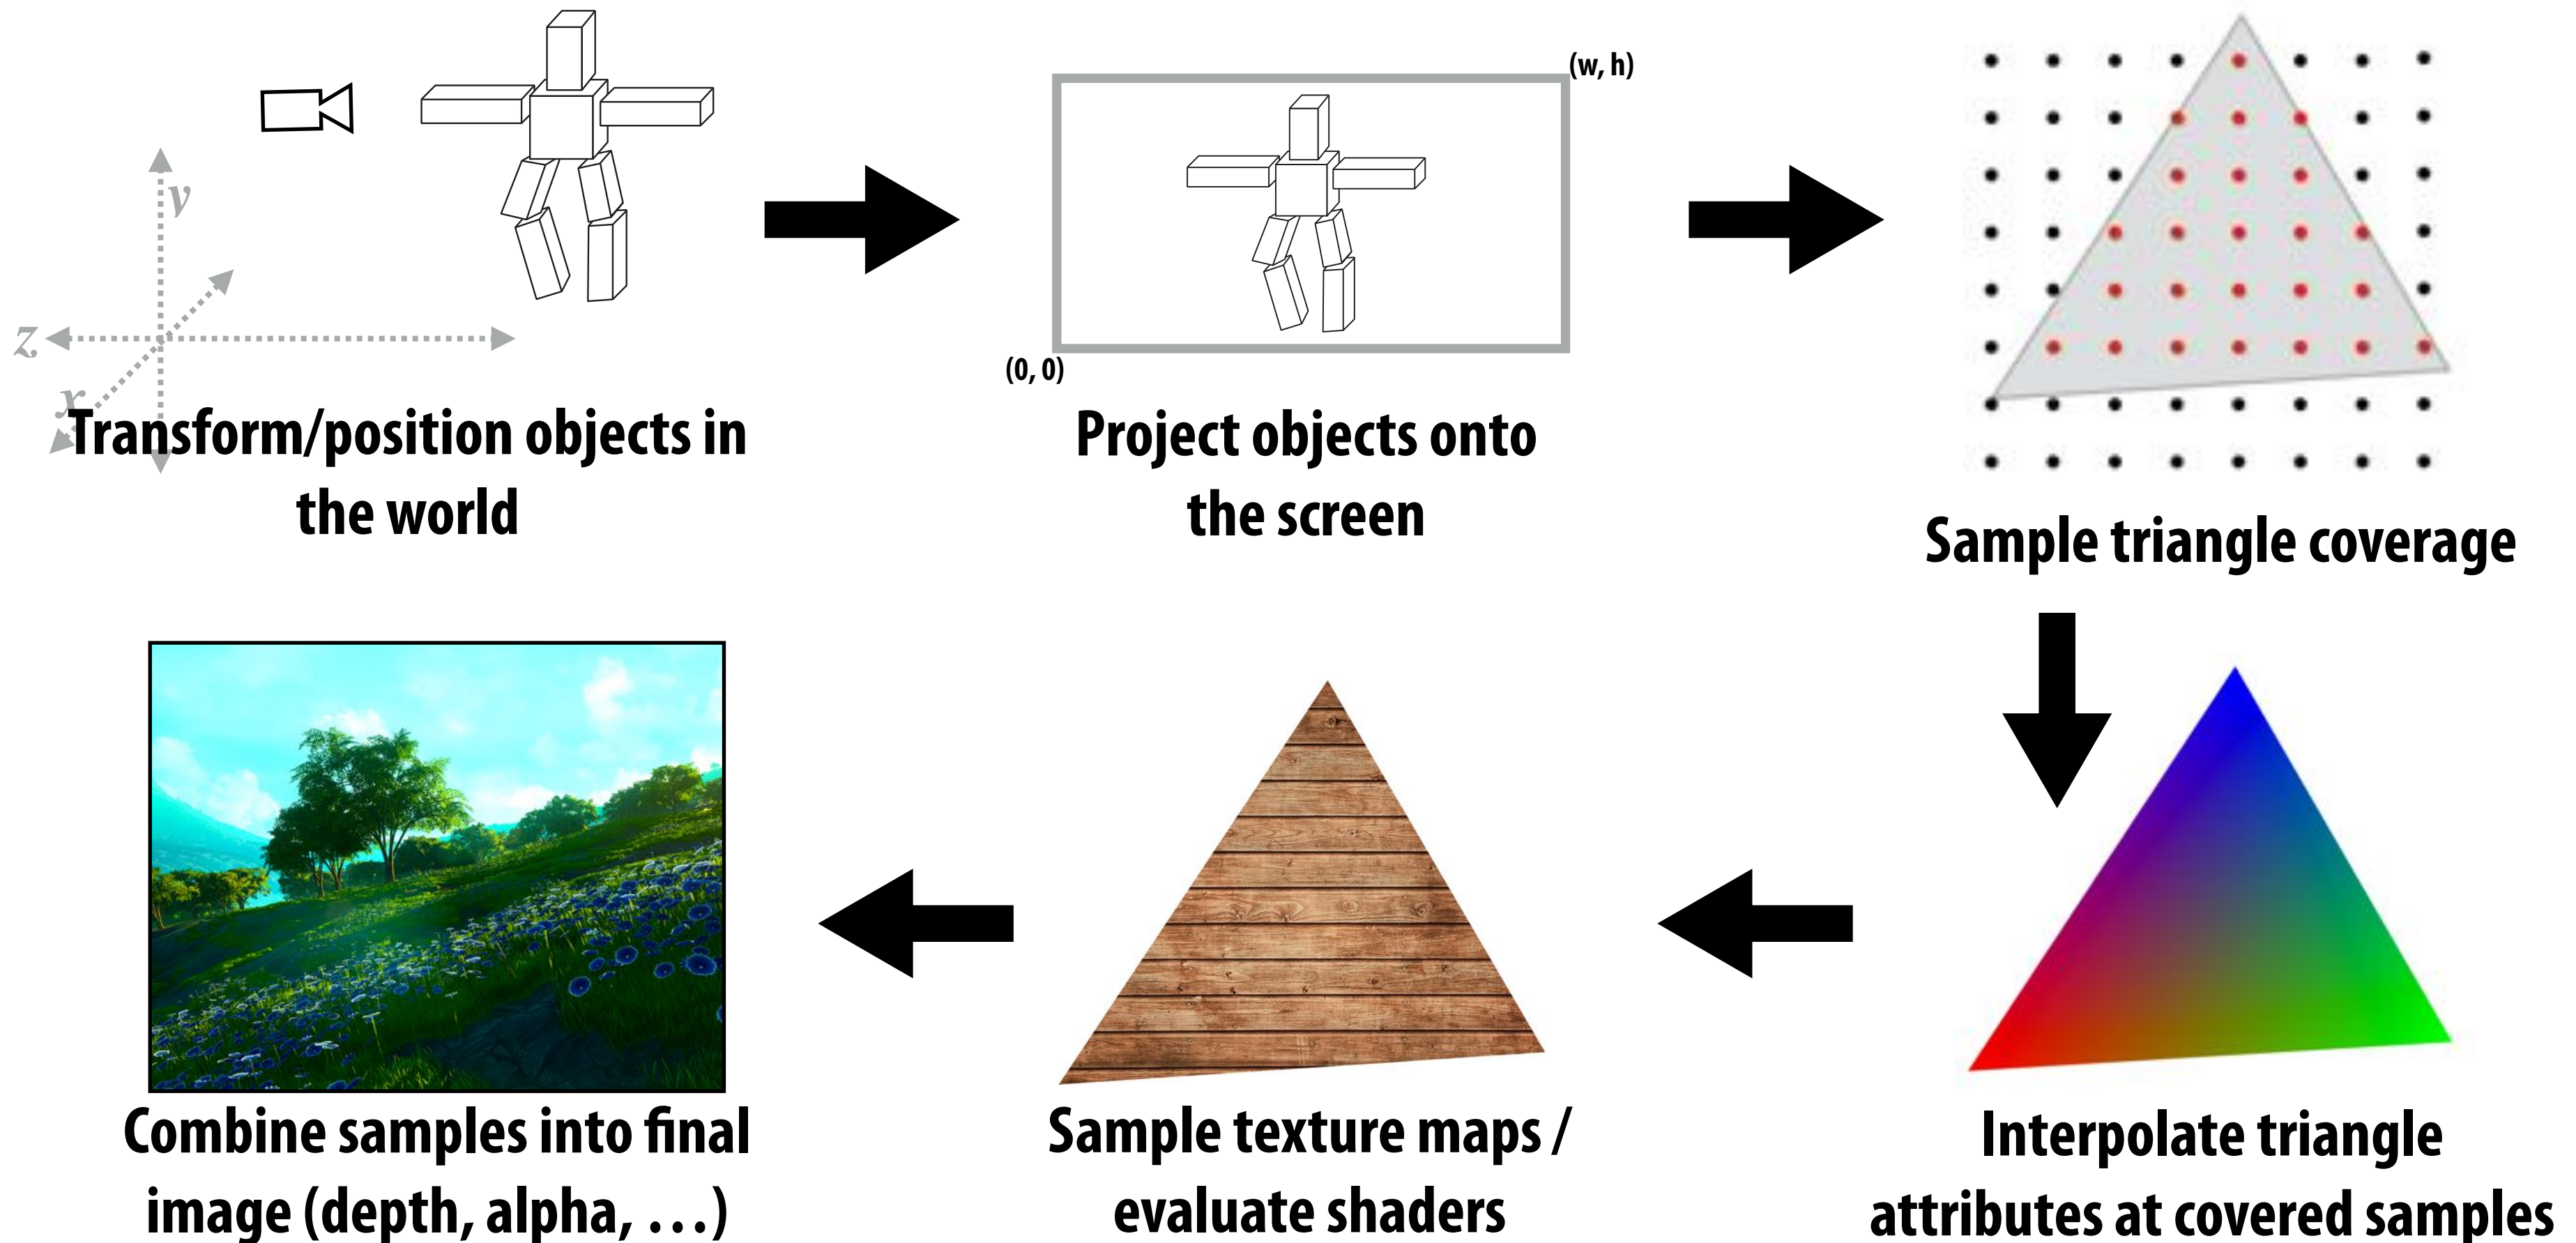
“point”



“line”

# The Rasterization Pipeline

Rough sketch of rasterization pipeline:




- Reflects standard “real world” pipeline (OpenGL/Direct3D)
  - the rest is just details (e.g., API calls)



# Let's draw some triangles on the screen


**Question 1: what pixels does the triangle overlap?  
("coverage")**



**Question 2: which triangle is closest to  
the camera in each pixel? ("occlusion")**


# The visibility problem

Recall the pinhole camera...



# The visibility problem

Recall the pinhole camera... which we can simplify with a “virtual sensor”:



## ■ Visibility problem in terms of rays:

- **COVERAGE:** What scene geometry is hit by a ray from a pixel through the pinhole?
- **OCCCLUSION:** Which object is the first hit along that ray?




# Computing triangle coverage

**“Which pixels does the triangle overlap?”**


**Input:**

**projected position of triangle vertices:  $P_0, P_1, P_2$**




**Output:**

**set of pixels “covered” by the triangle**




# What does it mean for a pixel to be covered by a triangle?

**Q: Which triangles “cover” this pixel?**




**One option: compute fraction of pixel area covered by triangle, then color pixel according to this fraction.**






# Coverage gets tricky when considering occlusion



**Pixel covered by triangle 1, other half covered by triangle 2**



**Interpenetration of triangles: even trickier**




**Two regions of triangle 1 contribute to pixel.  
One of these regions is not even convex.**

# Coverage via sampling

- Real scenes are complicated!
  - occlusion, transparency, ...
  - will talk about this more in a future lecture!
- Computing exact coverage is not practical
- Instead: view coverage as a sampling problem
  - don't compute exact/analytical answer
  - instead, test a collection of sample points
  - with enough points & smart choice of sample locations, can start to get a good estimate
- First, let's talk about sampling in general...




# Sampling 101: Sampling a 1D signal





# Sampling = taking measurements of a signal

Below: 5 measurements ("samples") of  $f(x)$




# Audio file: stores samples of a 1D signal



(most consumer audio is sampled 44,100 times per second, i.e., at 44.1 KHz)


**Reconstruction: given a set of samples, how might we attempt to reconstruct the original signal  $f(x)$ ?**






# Piecewise constant approximation

$\hat{f}(x)$  = value of sample closest to  $x$




# Piecewise linear approximation


$\hat{f}(x)$  = linear interpolation between values of two closest samples to  $x$



# How can we represent the signal more accurately?



# Reconstruction from denser sampling



..... = reconstruction via nearest

..... = reconstruction via linear interpolation



# 2D Sampling & Reconstruction


- Basic story doesn't change much for images:
  - sample values measure image (i.e., signal) at sample points
  - apply interpolation/reconstruction filter to approximate image



**original**



**piecewise constant  
("nearest neighbor")**





**piecewise bi-linear**




# Sampling 101: Summary

- **Sampling = measurement of a signal**
  - Encode signal as discrete set of samples
  - In principle, represent values at specific points (though hard to measure in reality!)
- **Reconstruction = generating signal from a discrete set of samples**
  - Construct a function that interpolates or approximates function values
  - E.g., piecewise constant/"nearest neighbor", or piecewise linear
  - Many more possibilities! For all kinds of signals (audio, images, geometry...)




# For rasterization, what function are we sampling?

$$\text{coverage}(x, y) := \begin{cases} 1, & \text{triangle contains point } (x, y) \\ 0, & \text{otherwise} \end{cases}$$




# Simple rasterization: just sample the coverage function



**Example:**  
Here I chose the coverage  
sample point to be at a  
point corresponding to the  
pixel center.




= triangle covers sample



= triangle does not cover sample

# Edge cases (literally)


**Is this sample point covered by triangle 1? or triangle 2? or both?**






# Breaking Ties\*


- When edge falls directly on a screen sample point, the sample is classified as within triangle if the edge is a “top edge” or “left edge”
  - Top edge: horizontal edge that is above all other edges
  - Left edge: an edge that is not exactly horizontal and is on the left side of the triangle. (triangle can have one or two left edges)



# Results of sampling triangle coverage




# I have a sampled signal, now I want to display it on a screen



# Pixels on a screen


**Each image sample sent to the display is converted into a little square of light of the appropriate color:  
(a pixel = picture element)**



**LCD display  
pixel on my  
laptop**

**\* Thinking of each LCD pixel as emitting a square of uniform intensity light of a single color is a bit of an approximation to how real displays work, but it will do for now.**


# So if we send the display this:





# We see this when we look at the screen

(assuming a screen pixel emits a square of perfectly uniform intensity of light)




# But the real coverage signal looked like this!



# Aliasing

# Sampling & Reconstruction

continuous signal  
(original)




**sample**

```
(-0.0457447  
-0.0209434  
0.0328632  
0.0468068  
0.0141212  
0.00506562  
0.0829571  
0.235369  
0.405682  
0.569204  
0.742511  
0.916946  
1.02  
0.973226  
0.781528  
0.539974  
0.3464  
0.223501  
0.134011  
0.0523279  
-0.00416744  
-0.0131345  
0.00959633  
0.0228676  
0.00789505  
)
```

**digital information**

continuous signal  
(approximate)




**reconstruct**


**Goal: reproduce original signal as accurately as possible.**

# 1D signal can be expressed as a superposition of frequencies


$$f_1(x) = \sin(\pi x)$$




$$f_2(x) = \sin(2\pi x)$$



$$f_4(x) = \sin(4\pi x)$$




$$f(x) = f_1(x) + 0.75 f_2(x) + 0.5 f_4(x)$$





# E.g., audio spectrum analyzer shows the amplitude of each frequency




↑  
Intensity of  
low-frequencies (bass)

↑  
Intensity of  
high frequencies

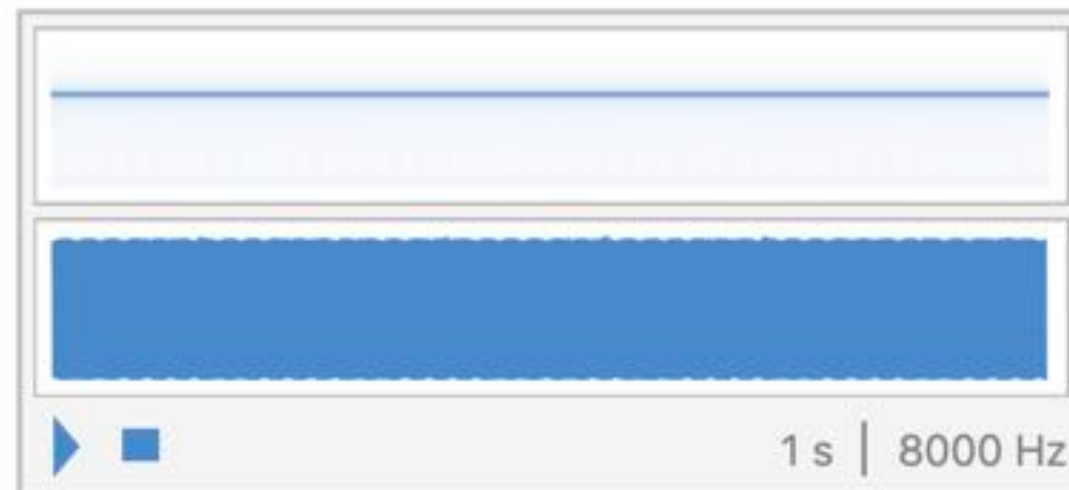
# Aliasing in Audio

Get a constant tone by playing a sinusoid of frequency  $\omega$ :


```
Play[Sin[4000 t], {t, 0, 1}]
```




```
Play[Sin[5000 t], {t, 0, 1}]
```




```
Play[Sin[6000 t], {t, 0, 1}]
```



Q: What happens if we increase  $\omega$  over time?




```
Play[Sin[ $\omega$  t], {t, 0, 10}]
```



Why did that happen?

# Undersampling high-frequency signals results in aliasing




**“Aliasing”: high frequencies in the original signal masquerade as low frequencies after reconstruction (due to undersampling)**

# Images can also be decomposed into “frequencies”



**Spatial domain result**




**Spectrum**



# Low frequencies only (smooth gradients)



**Spatial domain result**




**Spectrum (after low-pass filter)**  
All frequencies above cutoff have 0 magnitude

# Mid-range frequencies



**Spatial domain result**




**Spectrum (after band-pass filter)**



# Mid-range frequencies



**Spatial domain result**




**Spectrum (after band-pass filter)**

# High frequencies (edges)



**Spatial domain result  
(strongest edges)**



**Spectrum (after high-pass filter)  
All frequencies below threshold  
have 0 magnitude**


# An image as a sum of its frequency components




+



+



+




=






# Spatial aliasing: the function $\sin(x^2 + y^2)$





# Temporal aliasing: wagon wheel effect



**Camera's frame rate (temporal sampling rate) is too low for rapidly spinning wheel.**




# Nyquist-Shannon theorem

- Consider a band-limited signal: has no frequencies above some threshold  $\omega_0$ 
  - 1D example: low-pass filtered audio signal
  - 2D example: blurred image example from a few slides ago



- The signal can be perfectly reconstructed if sampled with period  $T = 1 / 2\omega_0$
- ...and if interpolation is performed using a “sinc filter”
  - ideal filter with no frequencies above cutoff (infinite extent!)








# Challenges of sampling in computer graphics

- **Signals are often not band-limited in computer graphics.**  
**Why?**

Hint:



- **Also, infinite extent of “ideal” reconstruction filter (sinc) is impractical for efficient implementations. Why?**





# Aliasing artifacts in images

- **Imperfect sampling + imperfect reconstruction leads to image artifacts**
  - **“Jaggies” in a static image**
  - **“Roping” or “shimmering” of images when animated**
  - **Moiré patterns in high-frequency areas of images**




# How can we reduce aliasing?

- No matter what we do, aliasing is a fact of life: any sampled representation eventually fails to capture frequencies that are too high.
- But we can still do our best to try to match sampling and reconstruction so that the signal we reproduce looks as much as possible like the signal we acquire
- For instance, if we think of a pixel as a “little square” of light, then we want the total light emitted to be the same as the total light in that pixel
  - I.e., we want to integrate the signal over the pixel (“box filter”)




Let's (approximately) integrate the signal **coverage (x,y)** by sampling...


# Initial coverage sampling rate (1 sample per pixel)



# Increase frequency of sampling coverage signal



# Supersampling







# Resampling


Converting from one discrete sampled representation to another





Original signal  
(high frequency edge)



Dense sampling of  
reconstructed signal




Reconstructed signal  
(lacks high frequencies)




Coarsely sampled signal

# Resample to display's pixel resolution


(Because a screen displays one sample value per screen pixel...)




# Resample to display's pixel rate (box filter)



# Resample to display's pixel rate (box filter)



# Displayed result (note anti-aliased edges)



# Recall: the real coverage signal was this






# Single Sample vs. Supersampling



**single sampling**




**2x2 supersampling**

# Single Sample vs. Supersampling



**single sampling**




**4x4 supersampling**



# Single Sample vs. Supersampling



**single sampling**




**32x32 supersampling**



# Checkerboard — Exact Solution

In very special cases we can compute the exact coverage:




Such cases are extremely rare—want solutions  
that will work in the general case!

**How do we actually evaluate  
 $\text{coverage}(x,y)$  for a triangle?**

# Point-in-triangle test

**Q: How do we check if a given point  $q$  is inside a triangle?**

**A: Check if it's contained in three half planes associated with the edges.**






# Point-in-triangle test

**Q: How do we check if a given point  $q$  is inside a triangle?**


**A: Check if it's contained in three half planes associated with the edges.**



# Point-in-triangle test

**Q: How do we check if a given point  $q$  is inside a triangle?**


**A: Check if it's contained in three half planes associated with the edges.**



# Point-in-triangle test

**Q: How do we check if a given point  $q$  is inside a triangle?**

**A: Check if it's contained in three half planes associated with the edges.**




# Point-in-triangle test

**Q:** How do we check if a given point  $q$  is inside a triangle?

**A:** Check if it's contained in three half planes associated with the edges.

Half plane test is then an exercise in linear algebra/vector calculus:



**GIVEN:** points  $P_i$ ,  $P_j$  along an edge, and a query point  $q$


**FIND:** whether  $q$  is to the “left” or “right” of the line from  $P_i$  to  $P_j$

(Careful to consider triangle coverage edge rules...)

# Traditional approach: incremental traversal

Since half-plane check looks very similar for different points, can save arithmetic by clever “incremental” schemes.


Incremental approach also visits pixels in an order that improves memory coherence: backtrack, zig-zag, Hilbert/Morton curves, ...






# Modern approach: parallel coverage tests

- Incremental traversal is very serial; modern hardware is highly parallel
- Alternative: test all samples in triangle “bounding box” in parallel
- Wide parallel execution overcomes cost of extra tests (most triangles cover many samples, especially when super-sampling)
- All tests share some “setup” calculations
- Modern graphics processing unit (GPU) has special-purpose hardware for efficiently performing point-in-triangle tests



**Q: What's a case where the naïve parallel approach is still very inefficient?**


# Naïve approach can be (very) wasteful...



# Hybrid approach: tiled triangle traversal


Idea: work “coarse to fine”:

- First, check if large blocks intersect the triangle
- If not, skip this block entirely (“early out”)
- If the block is contained inside the triangle, know all samples are covered (“early in”)
- Otherwise, test individual sample points in the block, in parallel




**This how real graphics hardware works!**

# Can we do even better for this example?



# Hierarchical strategies in computer graphics



**Q: Better way to find finest blocks?    A: Maybe: incremental traversal!**



# Summary

- Can frame many graphics problems in terms of sampling and reconstruction
  - sampling: turn a **continuous** signal into **digital** information
  - reconstruction: turn **digital** information into a **continuous** signal
  - aliasing occurs when the reconstructed signal presents a false sense of what the original signal looked like
- Can frame rasterization as sampling problem
  - sample coverage function into pixel grid
  - reconstruct by emitting a “little square” of light for each pixel
  - aliasing manifests as jagged edges, shimmering artifacts, ...
  - reduce aliasing via supersampling
- Triangle rasterization is basic building block for graphics pipeline
  - amounts to three half-plane tests
  - atomic operation—make it fast!
  - several strategies: incremental, parallel, blockwise, hierarchical...

# Next time: 3D Transformations

